

COMP9311 Week 9 Lecture

Relational Design Theory

Relational Design Theory

2/55

Previously we studied ER design and ER-to-relational mapping.

We claimed that this allows us to produce "good" schemas.

Can we make a stronger/formal statement on what makes a schema *good*?

The study of relational design theory

- examines some foundational notions of "schema goodness"
 - provides methods to transform schemas to make them better
-

... Relational Design Theory

3/55

Functional dependencies

- are constraints between attributes within a relation
- have implications for "good" relational schema design

What we study here:

- basic theory and definition of *functional dependencies*
- methodology for improving schema designs (*normalisation*)

The aim of studying this:

- improve understanding of relationships among data
 - gain enough formalism to assist practical database design
-

Relational Design and Redundancy

A *good* relational database design:

- must capture *all* necessary attributes/associations
- do this with a *minimal* amount of stored information

Minimal stored information \Rightarrow no redundant data.

In database design, *redundancy* is generally a "bad thing":

- causes problems maintaining consistency after updates

But ... redundancy may give performance improvements

- e.g. avoid a join to collect bits of data together

... Relational Design and Redundancy

5/55

Consider the following relation defining bank accounts/branches:

accountNo	balance	customer	branch	address	assets
A-101	500	1313131	Downtown	Brooklyn	9000000
A-102	400	1313131	Perryridge	Horseneck	1700000
A-113	600	9876543	Round Hill	Horseneck	8000000
A-201	900	9876543	Brighton	Brooklyn	7100000
A-215	700	1111111	Mianus	Horseneck	400000
A-222	700	1111111	Redwood	Palo Alto	2100000
A-305	350	1234567	Round Hill	Horseneck	8000000

We need to be careful updating this data, otherwise we may introduce inconsistencies.

... Relational Design and Redundancy

6/55

Insertion anomaly:

- when we insert a new record, we need to check that branch data is consistent with existing tuples

Update anomaly:

- if a branch changes address, we need to update all tuples referring to that branch

Deletion anomaly:

- if we remove information about the last account at a branch, all of the branch information disappears

(Insertion/update anomalies are avoidable, but need extra DBMS work)

... Relational Design and Redundancy

7/55

Insertion anomaly example (insert account A-306 at Round Hill):

accountNo	balance	customer	branch	address	assets
A-101	500	1313131	Downtown	Brooklyn	9000000
A-102	400	1313131	Perryridge	Horseneck	1700000
A-113	600	9876543	Round Hill	Horseneck	8000000
A-201	900	9876543	Brighton	Brooklyn	7100000
A-215	700	1111111	Mianus	Horseneck	400000
A-222	700	1111111	Redwood	Palo Alto	2100000
A-305	350	1234567	Round Hill	Horseneck	8000000
A-306	800	1111111	Round Hill	Horseneck	8000800

... Relational Design and Redundancy

8/55

Update anomaly example (update Round Hill branch address):

accountNo	balance	customer	branch	address	assets

A-101	500	1313131	Downtown	Brooklyn	9000000
A-102	400	1313131	Perryridge	Horseneck	1700000
A-113	600	9876543	Round Hill	<i>Palo Alto</i>	8000000
A-201	900	9876543	Brighton	Brooklyn	7100000
A-215	700	1111111	Mianus	Horseneck	400000
A-222	700	1111111	Redwood	Palo Alto	2100000
A-305	350	1234567	Round Hill	<i>Horseneck</i>	8000000

... Relational Design and Redundancy

9/55

Deletion anomaly example (remove account A-101):

accountNo	balance	customer	branch	address	assets
<i>A-101</i>	<i>500</i>	<i>1313131</i>	<i>Downtown</i>	<i>Brooklyn</i>	<i>9000000</i>
A-102	400	1313131	Perryridge	Horseneck	1700000
A-113	600	9876543	Round Hill	Horseneck	8000000
A-201	900	9876543	Brighton	Brooklyn	7100000
A-215	700	1111111	Mianus	Horseneck	400000
A-222	700	1111111	Redwood	Palo Alto	2100000
A-305	350	1234567	Round Hill	Horseneck	8000000

Where is the Downtown branch located? What are its assets?

Database Design (revisited)

10/55

To avoid these kinds of update problems:

- need a schema with "minimal overlap" between tables
- each table contains a "coherent" collection of data values

One way to achieve this:

- start with a "universal relation" U (all relevant attrs)
- *decompose* relation U into several smaller relations R_i
- where each R_i has "minimal overlap" with other R_j
- but there is sufficient overlap to reconstruct original table

Typically, each R_i contains info about one entity (e.g. branch, customer, ...)

... Database Design (revisited)

11/55

But we already have a design procedure (ER-then-relational)

- and it appears to produce schemas without redundancy

Why do we need another design procedure?

1. ER design does not guarantee minimal redundancy
 - dependency theory allows us to check designs for residual problems
2. the new procedure gives (semi)automated design
 - determine all of the attributes in the problem domain
 - collect them all together in a "universal relation"
 - provide information about how attributes are related
 - apply *normalisation* to decompose into non-redundant relations

Notation/Terminology

12/55

Most texts adopt the following terminology:

Relation schemas	upper-case letters, denoting set of all attributes (e.g. R, S, P, Q)
Relation instances	lower-case letter corresponding to schema (e.g. $r(R), s(S), p(P), q(Q)$)
Tuples	lower-case letters (e.g. t, t', t_I, u, v)
Attributes	upper-case letters from start of alphabet (e.g. A, B, C, D)
Sets of attributes	simple concatenation of attribute names (e.g. $X=ABCD, Y=EFG$)

Attributes in tuples tuple[attrSet] (e.g. $t[ABCD]$, $t[X]$)

Functional Dependency

13/55

A relation instance $r(R)$ satisfies a dependency $X \rightarrow Y$ if

- for any $t, u \in r$, $t[X] = u[X] \Rightarrow t[Y] = u[Y]$

In other words, if two tuples in R agree in their values for the set of attributes X , then they must also agree in their values for the set of attributes Y .

We say that " Y is *functionally dependent* on X ".

Attribute sets X and Y may overlap; trivially true that $X \rightarrow X$.

Notes:

- $X \rightarrow Y$ can also be read as " X determines Y "
- the single arrow \rightarrow denotes *functional dependency*
- the double arrow \Rightarrow denotes logical implication

... Functional Dependency

14/55

Consider the following instance $r(R)$ of the relation schema $R(ABCDE)$:

A	B	C	D	E
a_1	b_1	c_1	d_1	e_1
a_2	b_1	c_2	d_2	e_1
a_3	b_2	c_1	d_1	e_1
a_4	b_2	c_2	d_2	e_1
a_5	b_3	c_3	d_1	e_1

What kind of dependencies can we observe among the attributes in $r(R)$?

... Functional Dependency

Since A values are unique, the definition of fd gives:

- $A \rightarrow B, A \rightarrow C, A \rightarrow D, A \rightarrow E$
- $A \rightarrow BC, A \rightarrow CD, \dots A \rightarrow BCDE$
- can be summarised as $A \rightarrow BCDE$

Since all E values are the same, it follows that:

- $A \rightarrow E, B \rightarrow E, C \rightarrow E, D \rightarrow E$
- in general, *cannot* be summarised as $ABCD \rightarrow E$

... Functional Dependency

Other observations:

- combinations of BC are unique, therefore $BC \rightarrow ADE$
- combinations of BD are unique, therefore $BD \rightarrow ACE$
- if C values match, so do D values, therefore $C \rightarrow D$
- however, D values don't determine C values, so $\neg(D \rightarrow C)$

We could derive many other dependencies, e.g. $AE \rightarrow BC, \dots$

In practice, choose a minimal set of fds (*basis*)

- from which all other fds can be derived
- which captures useful problem-domain information

Exercise: Functional Dependencies (1)

Real estate agents conduct visits to rental properties

- need to record which property, who went, when, results
- each property is assigned a unique code (P#, e.g. PG4)
- each staff member has a staff number (S#, e.g. SG43)
- staff members use company cars to conduct visits

- a visit occurs at a specific time on a given day
- notes are made on the state of the property after each visit

The company stores all of the associated data in a spreadsheet.

... Exercise: Functional Dependencies (1)

18/55

Describe any functional dependencies that exist in this data:

P#	When	Address	Notes	S#	Name	CarReg
PG4	03/06 15:15	55 High St	Bathroom leak	SG44	Rob	ABK754
PG1	04/06 11:10	47 High St	All ok	SG44	Rob	ABK754
PG4	03/07 12:30	55 High St	All ok	SG43	Dave	ATS123
PG1	05/07 15:00	47 High St	Broken window	SG44	Rob	ABK754
PG1	05/07 15:00	47 High St	Leaking tap	SG44	Rob	ABK754
PG2	13/07 12:00	12 High St	All ok	SG42	Peter	ATS123
PG1	10/08 09:00	47 High St	Window fixed	SG42	Peter	ATS123
PG3	11/08 14:00	99 High St	All ok	SG41	John	AAA001
PG4	13/08 10:00	55 High St	All ok	SG44	Rob	ABK754
PG3	05/09 11:15	99 High St	Bathroom leak	SG42	Peter	ATS123

State assumptions used in determining the *fds*.

Exercise: Functional Dependencies (2)

19/55

What functional dependencies exist in the following table:

A	B	C	D
1	a	6	x
2	b	7	y
3	c	7	z
4	d	6	x
5	a	6	y
6	b	7	z
7	c	7	x
8	d	6	y

How is this case different to the previous one?

Functional Dependency

20/55

The above discussion was about dependency within a relation instance $r(R)$.

Much more important for design is the notion of dependency across all possible instances of the relation (i.e. a schema-based dependency).

This is a simple generalisation of the previous definition:

- for any $t, u \in \text{any } r(R)$, $t[X] = u[X] \Rightarrow t[Y] = u[Y]$

Useful because such dependencies reflect the semantics of the problem domain.

... Functional Dependency

21/55

Can we generalise some ideas about functional dependency?

E.g. are there dependencies that hold for *any* relation?

- yes, but they're generally trivial, e.g. $Y \subset X \Rightarrow X \rightarrow Y$

E.g. do some dependencies suggest the existence of others?

- yes, *rules of inference* allow us to *derive* dependencies
 - allows us to reason about sets of functional dependencies
-

Inference Rules

22/55

Armstrong's rules are complete, general rules of inference on *fds*.

F1. *Reflexivity* e.g. $X \rightarrow X$

- a formal statement of *trivial dependencies*; useful for derivations

F2. *Augmentation* e.g. $X \rightarrow Y \Rightarrow XZ \rightarrow YZ$

- if a dependency holds, then we can freely expand its left hand side

F3. *Transitivity* e.g. $X \rightarrow Y, Y \rightarrow Z \Rightarrow X \rightarrow Z$

- the "most powerful" inference rule; useful in multi-step derivations

... Inference Rules

23/55

Armstrong's rules are complete, but other useful rules exist:

F4. *Additivity* e.g. $X \rightarrow Y, X \rightarrow Z \Rightarrow X \rightarrow YZ$

- useful for constructing new right hand sides of *fds* (also called *union*)

F5. *Projectivity* e.g. $X \rightarrow YZ \Rightarrow X \rightarrow Y, X \rightarrow Z$

- useful for reducing right hand sides of *fds* (also called *decomposition*)

F6. *Pseudotransitivity* e.g. $X \rightarrow Y, YZ \rightarrow W \Rightarrow XZ \rightarrow W$

- shorthand for a common transitivity derivation

... Inference Rules

24/55

Example: determining validity of $AB \rightarrow GH$, given:

$$R = ABCDEFGHIJ$$

$$F = \{ AB \rightarrow E, AG \rightarrow J, BE \rightarrow I, E \rightarrow G, GI \rightarrow H \}$$

Derivation:

1. $AB \rightarrow E$ (given)
2. $AB \rightarrow AB$ (using F1)
3. $AB \rightarrow B$ (using F5 on 2)
4. $AB \rightarrow BE$ (using F4 on 1,3)
5. $BE \rightarrow I$ (given)

6. $AB \rightarrow I$ (using F3 on 4,5)
7. $E \rightarrow G$ (given)
8. $AB \rightarrow G$ (using F3 on 1,7)
9. $AB \rightarrow GI$ (using F4 on 6,8)
10. $GI \rightarrow H$ (given)
11. $AB \rightarrow H$ (using F3 on 6,8)
12. $AB \rightarrow GH$ (using F4 on 8,11)

Closures

25/55

Given a set F of fds , how many new fds can we derive?

For a finite set of attributes, there must be a finite set of fds .

The largest collection of dependencies that can be derived from F is called the *closure* of F and is denoted F^+ .

Closures allow us to answer two interesting questions:

- is a particular dependency $X \rightarrow Y$ derivable from F ?
- are two sets of dependencies F and G equivalent?

... Closures

26/55

For the question "is $X \rightarrow Y$ derivable from F ?" ...

- compute the closure F^+ ; check whether $X \rightarrow Y \in F^+$

For the question "are F and G equivalent?" ...

- compute closures F^+ and G^+ ; check whether they're equal

Unfortunately, closures can be very large, e.g.

$R = ABC, \quad F = \{ AB \rightarrow C, \quad C \rightarrow B \}$

$F^+ = \{ A \rightarrow A, \quad AB \rightarrow A, \quad AC \rightarrow A, \quad AB \rightarrow B, \quad BC \rightarrow B, \quad ABC \rightarrow B, \\ C \rightarrow C, \quad AC \rightarrow C, \quad BC \rightarrow C, \quad ABC \rightarrow C, \quad AB \rightarrow AB, \quad \dots, \\ AB \rightarrow ABC, \quad AB \rightarrow ABC, \quad C \rightarrow B, \quad C \rightarrow BC, \quad AC \rightarrow B, \quad AC \rightarrow AB \}$

... Closures

27/55

Algorithms based on F^+ rapidly become infeasible.

To solve this problem, use closures based on sets of attributes rather than sets of *fds*.

Given a set X of attributes and a set F of *fds*, the largest set of attributes that can be derived from X using F , is called the *closure* of X (denoted X^+).

We can prove (using additivity) that $(X \rightarrow Y) \in F^+$ iff $Y \subset X^+$.

For computation, $|X^+|$ is bounded by the number of attributes.

... Closures

28/55

For the question "is $X \rightarrow Y$ derivable from F ?" ...

- compute the closure X^+ , check whether $Y \subset X^+$

For the question "are F and G equivalent?" ...

- for each dependency in G , check whether derivable from F
- for each dependency in F , check whether derivable from G
- if true for all, then $F \Rightarrow G$ and $G \Rightarrow F$ which implies $F^+ = G^+$

For the question "what are the keys of R implied by F ?" ...

- find subsets $K \subset R$ such that $K^+ = R$

Exercise: Computing Keys

29/55

Determine possible primary keys for each of the following:

1. $FD = \{ A \rightarrow B, C \rightarrow D, E \rightarrow FG \}$
2. $FD = \{ A \rightarrow B, B \rightarrow C, C \rightarrow D \}$
3. $FD = \{ A \rightarrow B, B \rightarrow C, C \rightarrow A \}$
4. $FD = \{ ABH \rightarrow C, A \rightarrow D, C \rightarrow E, F \rightarrow A, E \rightarrow F, BGH \rightarrow E \}$

Normalization

30/55

Normalization: branch of relational theory providing design insights.

The goals of normalization:

- be able to characterise the level of redundancy in a relational schema
- provide mechanisms for transforming schemas to remove redundancy

Normalization draws heavily on the theory of functional dependencies.

Normal Forms

31/55

Normalization theory defines six *normal forms* (NFs).

- First, Second, Third Normal Forms (1NF, 2NF, 3NF) (Codd 1972)
- Boyce-Codd Normal Form (BCNF) (1974)
- Fourth Normal Form (4NF) (Zaniolo 1976, Fagin 1977)
- Fifth Normal Form (5NF) (Fagin 1979)

NF hierarchy: $5NF \Rightarrow 4NF \Rightarrow BCNF \Rightarrow 3NF \Rightarrow 2NF \Rightarrow 1NF$

1NF allows most redundancy; 5NF allows least redundancy.

... Normal Forms

32/55

1NF

all attributes have atomic values
we assume this as part of relational model,

so *every* relation schema is in 1NF

2NF	all non-key attributes fully depend on key (i.e. no partial dependencies) avoids much redundancy
3NF BCNF	no attributes dependent on non-key attrs (i.e. no transitive dependencies) avoids most remaining redundancy

... Normal Forms

33/55

The normalization process:

- decide which normal form rNF is "acceptable".
 - how much redundancy are we willing to tolerate?
- check whether each relation in schema is in rNF
- if a relation is not in rNF
 - partition into sub-relations where each is closer to rNF
- repeat until all relations in schema are in rNF

... Normal Forms

34/55

In practice, BCNF and 3NF are the most important.
(these are the "acceptable normal forms" for relational design)

Boyce-Codd Normal Form (BCNF):

- eliminates all redundancy due to functional dependencies
- but may not preserve original functional dependencies

Third Normal Form (3NF):

- eliminates most (but not all) redundancy due to fd s
- guaranteed to preserve all functional dependencies

35/55

Relation Decomposition

The standard transformation technique to remove redundancy:

- *decompose* relation R into relations S and T

We accomplish decomposition by

- selecting (overlapping) subsets of attributes
- forming new relations based on attribute subsets

Properties: $R = S \cup T$, $S \cap T \neq \emptyset$ and $r(R) = s(S) \bowtie t(T)$

May require several decompositions to achieve acceptable NF.

Normalization algorithms tell us how to choose S and T .

Schema Design

36/55

Consider the following relation for *BankLoans*:

branchName	branchCity	assets	custName	loanNo	amount
Downtown	Brooklyn	9000000	Jones	L-17	1000
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
Downtown	Brooklyn	9000000	Jackson	L-15	1500
Mianus	Horseneck	400000	Jones	L-93	500
Round Hill	Horseneck	8000000	Turner	L-11	900
North Town	Rye	3700000	Hayes	L-16	1300

This schema has all of the update anomalies mentioned earlier.

... Schema Design

37/55

To improve the design, decompose the *BankLoans* relation.

The following decomposition is not helpful:

Branch(branchName, branchCity, assets)

CustLoan(custName, loanNo, amount)

because we lose information (which branch is a loan held at?)

Another possible decomposition:

BranchCust(branchName, branchCity, assets, custName)

CustLoan(custName, loanNo, amount)

... Schema Design

38/55

The *BranchCust* relation instance:

branchName	branchCity	assets	custName
Downtown	Brooklyn	9000000	Jones
Redwood	Palo Alto	2100000	Smith
Perryridge	Horseneck	1700000	Hayes
Downtown	Brooklyn	9000000	Jackson
Mianus	Horseneck	400000	Jones
Round Hill	Horseneck	8000000	Turner
North Town	Rye	3700000	Hayes

... Schema Design

39/55

The *CustLoan* relation instance:

custName	loanNo	amount
Jones	L-17	1000

Smith	L-23	2000
Hayes	L-15	1500
Jackson	L-15	1500
Jones	L-93	500
Turner	L-11	900
Hayes	L-16	1300

... Schema Design

40/55

Now consider the result of (*BranchCust Join CustLoan*)

branchName	branchCity	assets	custName	loanNo	amount
Downtown	Brooklyn	9000000	Jones	L-17	1000
<i>Downtown</i>	<i>Brooklyn</i>	<i>9000000</i>	<i>Jones</i>	<i>L-93</i>	<i>500</i>
Redwood	Palo Alto	2100000	Smith	L-23	2000
Perryridge	Horseneck	1700000	Hayes	L-15	1500
<i>Perryridge</i>	<i>Horseneck</i>	<i>1700000</i>	<i>Hayes</i>	<i>L-16</i>	<i>1300</i>
Downtown	Brooklyn	9000000	Jackson	L-15	1500
Mianus	Horseneck	400000	Jones	L-93	500
<i>Mianus</i>	<i>Horseneck</i>	<i>400000</i>	<i>Jones</i>	<i>L-17</i>	<i>1000</i>
Round Hill	Horseneck	8000000	Turner	L-11	900
North Town	Rye	3700000	Hayes	L-16	1300
<i>North Town</i>	<i>Rye</i>	<i>3700000</i>	<i>Hayes</i>	<i>L-15</i>	<i>1500</i>

... Schema Design

41/55

This is clearly not a successful decomposition.

The fact that we ended up with extra tuples was symptomatic of losing some critical "connection" information during the decomposition.

Such a decomposition is called a *lossy decomposition*.

In a good decomposition, we should be able to reconstruct the original relation exactly:

$$\text{if } R \text{ is decomposed into } S \text{ and } T, \text{ then } \text{Join}(S, T) = R$$

Such a decomposition is called *lossless join decomposition*.

Boyce-Codd Normal Form

42/55

A relation schema R is in BCNF w.r.t a set F of functional dependencies iff:

for all $fds X \rightarrow Y$ in F^+

- either $X \rightarrow Y$ is trivial (i.e. $Y \subset X$)
- or X is a superkey

A DB schema is in BCNF if all relation schemas are in BCNF.

Observations:

- any two-attribute relation is in BCNF
 - any relation with key K , other attributes X , and $K \rightarrow X$ is in BCNF
-

... Boyce-Codd Normal Form

43/55

If we transform a schema into BCNF, we are guaranteed:

- no update anomalies due to *fd*-based redundancy
- lossless join decomposition

However, we are *not* guaranteed:

- the new schema preserves all *fds* from the original schema

This may be a problem if the *fds* contain significant semantic information about the problem domain.

If we need to preserve dependencies, use 3NF.

BCNF Decomposition

44/55

The following algorithm converts an arbitrary schema to BCNF:

Inputs: schema R , set F of fd s
Output: set Res of BCNF schemas

```
Res = {R};  
while (any schema  $S \in Res$  is not in BCNF) {  
    choose any  $fd\ X \rightarrow Y$  on  $S$  that violates BCNF  
     $Res = (Res - S) \cup (S - Y) \cup XY$   
}
```

The last step means:

- make a table from XY ; drop Y from table S

... BCNF Decomposition

45/55

Example (the *BankLoans* schema):

BankLoans(*branchName*, *branchCity*, *assets*, *custName*, *loanNo*, *amount*)

Has functional dependencies F

- $branchName \rightarrow assets, branchCity$
- $loanNo \rightarrow amount, branchName$

The key for *BankLoans* is *branchName*, *custName*, *loanNo*

... BCNF Decomposition

46/55

Applying the BCNF algorithm:

- check *BankLoans* relation ... it is not in BCNF
($branchName \rightarrow assets, branchCity$ violates BCNF criteria; LHS is not a key)

- to fix ... decompose *BankLoans* into

Branch(branchName, branchCity, assets)

LoanInfo(branchName, custName, loanNo, amount)

- check *Branch* relation ... it is in BCNF
(the only nontrivial *fds* have LHS=*branchName*, which is a key)

(continued)

... BCNF Decomposition

47/55

Applying the BCNF algorithm (cont):

- check *LoanInfo* relation ... it is not in BCNF
(*loanNo* \rightarrow *amount, branchName* violates BCNF criteria; LHS is not a key)
- to fix ... decompose *LoanInfo* into

Loan(branchName, loanNo, amount)

Borrower(custName, loanNo)

- check *Loan* ... it is in BCNF
- check *Borrower* ... it is in BCNF

Exercise: BCNF Decomposition (1)

48/55

Consider the schema *R* and set of *fds* *F*

R = *ABCDEFGH*

F = { *ABH* \rightarrow *C*, *A* \rightarrow *DE*, *BGH* \rightarrow *F*, *F* \rightarrow *ADH*, *BH* \rightarrow *GE* }

Produce a BCNF decomposition of *R*.

Exercise: BCNF Decomposition (2)

49/55

Recall the following table (based on e.g. a spreadsheet):

P#	When	Address	Notes	S#	Name	CarReg
PG4	03/06 15:15	55 High St	Bathroom leak	SG44	Rob	ABK754
PG1	04/06 11:10	47 High St	All ok	SG44	Rob	ABK754
PG4	03/07 12:30	55 High St	All ok	SG43	Dave	ATS123
PG1	05/07 15:00	47 High St	Broken window	SG44	Rob	ABK754
PG1	05/07 15:00	47 High St	Leaking tap	SG44	Rob	ABK754
PG2	13/07 12:00	12 High St	All ok	SG42	Peter	ATS123
...						

Recall the functional dependencies identified previously,

Use these to convert the table to a BCNF schema.

Third Normal Form

50/55

A relation schema R is in 3NF w.r.t a set F of functional dependencies iff:

for all $fds X \rightarrow Y$ in F^+

- either $X \rightarrow Y$ is trivial (i.e. $Y \subset X$)
- or X is a superkey
- or Y is a single attribute from a key

A DB schema is in 3NF if all relation schemas are in 3NF.

The extra condition represents a slight weakening of BCNF requirements.

... Third Normal Form

51/55

If we transform a schema into 3NF, we are guaranteed:

- lossless join decomposition
- the new schema preserves all of the fds from the original schema

However, we are *not* guaranteed:

- no update anomalies due to *fd*-based redundancy

Whether to use BCNF or 3NF depends on overall design considerations.

... Third Normal Form

52/55

The following algorithm converts an arbitrary schema to 3NF:

Inputs: schema R , set F of *fds*

Output: set R_i of 3NF schemas

```

let  $F_C$  be a minimal cover for  $F$ 
 $Res = \{\}$ 
for each fd  $X \rightarrow Y$  in  $F_C$  {
    if (no schema  $S \in Res$  contains  $XY$ ) {
         $Res = Res \cup XY$ 
    }
}
if (no schema  $S \in Res$  contains a candidate key for  $R$ ) {
     $K =$  any candidate key for  $R$ 
     $Res = Res \cup K$ 
}

```

... Third Normal Form

53/55

Critical step is producing minimal cover F_C for F

A set F of *fds* is minimal if

- every *fd* $X \rightarrow Y$ is *simple*
(Y is a single attribute)
- every *fd* $X \rightarrow Y$ is *left-reduced*
(no $Z \subset X$ such that $Z \rightarrow A$ could replace $X \rightarrow A$ in F and preserve F^+)
- every *fd* $X \rightarrow Y$ is *necessary*
(no $X \rightarrow Y$ can be removed without changing F^+)

Algorithm: right-reduce, left-reduce, eliminate redundant *fds*

Exercise: 3NF Decomposition (1)

Consider the schema R and set of fds F

$R = ABCDEFGH$

$F = F_c = \{ ABH \rightarrow C, A \rightarrow D, C \rightarrow E, F \rightarrow A, E \rightarrow F, BGH \rightarrow E \}$

Produce a 3NF decomposition of R .

Database Design Methodology

To achieve a "good" database design:

- identify attributes, entities, relationships \rightarrow ER design
- map ER design to relational schema
- identify constraints (including keys and functional dependencies)
- apply BCNF/3NF algorithms to produce normalized schema

Note: may subsequently need to "denormalise" if the design yields inadequate performance.

Produced: 19 Sep 2012