COMP9311 Week 06 Lecture

Extending SQL

Limitations of Basic SQL

What we have seen of SQL so far:

- data definition language (`create table(...)`)
- constraints   (domain, key, referential integrity)
- query language (`select...from...where...`)
- views   (give names to SQL queries)

This is not sufficient to write complete applications.

Need to write *programs* to manipulate database.

Extending database functionality would also help.

Extending SQL

Ways in which standard SQL might be extended:

- new data types (incl. constraints, I/O, indexes, ...)
- more operations/aggregates for use in queries
- more powerful constraint checking
- event-based triggered actions
- different kinds of queries (e.g. recursive)

All are potentially useful in application development.

New Data Types

SQL data definition language provides:

- atomic types: integer, float, character, boolean
- ability to define tuple types (`create table`)

SQL also provides mechanisms to define new types, e.g.

```
create domain Positive as integer check (value > 0);
create type Rating as enum ('poor','ok','excellent');
create type Pair as (x integer, y integer);
```

`CREATE TYPE` also allows new basic types to be defined.

New Functions

SQL provides for new functions via *stored procedures*.

PostgreSQL has functions in: SQL, PLpgSQL, Python, ...

```
create function
   f(arg₁ type₁, arg₂ type₂, ...) returns type
as $$ ... function body ... $$
language language [ mode ];
```

Possible modes (can make *big* performance difference)

- `immutable` ... does not access database (fast)
- `stable` ... does not modify the database
- `volatile` ... may change the database (slow, default)

Exercise: functions on (sets of) integers

Write PLpgSQL functions:

```
-- factorial n!
function fac(n integer) returns integer

-- returns integers 1..hi
function iota(hi integer) returns setof integer
```

```
-- returns integers lo..hi
function iota(lo integer, hi integer)
    returns setof integer

-- returns integers lo,lo+inc,..hi
function iota(lo integer, hi integer, step integer)
    returns setof integer
```

Queries

Advanced Query Types

Many specialised types of query have been identified.

We have seen: select/project/join, aggregation, grouping

Many modern queries (e.g. skyline) come from OLAP.

Two important standard query types:

- *recursive* ... e.g. to manage hierarchies, graphs (1999)
- *window* ... e.g. to "spread" group-by summaries (2003)

Window Functions

Group-by allows us to

- summarize a set of tuples
- that have common values for a set of attributes

E.g. average mark for each student (not the WAM)

```
select student, avg(mark)
from   CourseEnrolments
```

```
group  by student;
```

Produces a single summary tuple for each group.

... Window Functions

Window functions allow us to

- compute summary values for a group
- append the summary value to each tuple in the group

E.g. attach student's average mark to each enrolment

```
select *, avg(mark)
over   (partition by student)
from   CourseEnrolments;
```

... Window Functions

Comparison of `group by` and `partition by`:

```
select student,avg(mark) ... group by student

 student  |  avg
----------+-------
 46000936 | 64.75
 46001128 | 73.50
```

```
select *,avg(mark) over (partition by student) ...

 student  | course | mark | grade | stueval |  avg
----------+--------+------+-------+---------+-------
 46000936 |  11971 |   68 | CR    |       3 | 64.75
 46000936 |  12937 |   63 | PS    |       3 | 64.75
 46000936 |  12045 |   71 | CR    |       4 | 64.75
 46000936 |  11507 |   57 | PS    |       2 | 64.75
 46001128 |  12932 |   73 | CR    |       3 | 73.50
 46001128 |  13498 |   74 | CR    |       5 | 73.50
```

Group by cannot give you the detail,
you need to use select *, ··· over (partition by xxx)

```
46001128 |   11909 |   79 | DN      |        4 | 73.50
46001128 |   12118 |   68 | CR      |        4 | 73.50
```

Exercise: Underachievers

Using window functions, write an SQL function to find:

- all of the students in a given course
- whose mark is < 60% of average mark for course

```
create or replace function
   under(integer) returns setof CourseEnrolments
as $$
...
$$ language sql;
```

WITH Queries

We often break a complex query up into views, e.g.

```
create view V as
select a,b,c from ... where ...;

create view W as
select d,e from ... where ...;

create view Result(x,y,z) as
select V.a, V.b, W.e
from   V join W on (v.c = W.d);
```
In this case, the v and w are called outside

But we don't always need/want the views to persist.

... WITH Queries

WITH allows scoped/temporary views, e.g.   *With is a local variable*

```
with V as (select a,b,c from ... where ...),
     W as (select d,e from ... where ...)
select V.a as x, V.b as y, W.e as z
from   V join W on (v.c = W.d);
```

The views `V` and `W`

- only exist while this query is evaluated
- are not accessible in any other context

`V` and `W` are also called "common table expressions" (CTEs)

... `WITH` Queries

Note that named subqueries achieve the same effect:

```
select V.a as x, V.b as y, W.e as z
from   (select a,b,c from ... where ...) as V,
       (select d,e from ... where ...) as W
where  V.c = W.d;
```
This way is easier to read

For this purpose, `WITH` is a syntactic convenience.

However, `WITH` also provides recursive queries.

Recursive Queries

Recursive queries are defined as: *with recursive -> this is the key word*

```
with recursive T(a₁, a₂, ...) as
(
    non-recursive select
  union
    recursive select involving T
)
select ... from T where ...
```

`T(a₁, a₂, ...)` is a recursively-defined view.

... Recursive Queries

Example: generate sum of first 100 integers

```
with recursive nums(n) as (
    select 1
  union
    select n+1 from nums where n < 100
)
select sum(n) from nums;
-- which produces ...
 sum
------
 5050
```

... Recursive Queries

The subqueries generating `T` cannot be arbirtrary

- non-recursive select
  - does not refer to `T`
  - generates an initial set of tuples for `T` *(initialisation)*
- recursive select
  - must be a query involving `T`
  - must include a `where` condition
  - must eventually return an empty result *(termination)*

... Recursive Queries

Semantics of recursive query yielding `T(a₁, a₂, ...)`:

```
-- res, work, tmp are all temporary tables
res = result of non-recursive query
work = res
```

Recursion in SQL is different from that in programming languages
It does not cost more memory
It will not crash the process, but it will also kill.

```
while (work is not empty) {
    -- using work as the value for T ...
    tmp = result of recursive query
    res = res + tmp
    work = tmp
}
return res
```

Not more than twice for Assignement 2

Exercise: Recursive Queries   *The answer will be post in the lab*

Re-write the `iota(lo,hi)` function so that it uses a recursive query.

```
function iotar(integer,integer) returns integer
```

Re-write the `facultyOf(ouid)` function so that it uses a recursive query.

```
function facOf(integer) returns integer
```

Aggregates

Aggregates

Aggregates reduce a collection of values into a single result.

Examples: `count(`*Tuples*`)`, `sum(`*Numbers*`)`, etc.

The action of an aggregate function can be viewed as:

```
AggState = initial state
for each item V {
    # update AggState to include V
    AggState = newState(AggState, V)
}
return final(AggState)
```

... Aggregates

Aggregates are commonly used with `GROUP BY`.

In that context, they "summarise" each group.

Example:

```
R                   select a,sum(b),count(*)
 a | b | c          from R group by a
---+---+---
 1 | 2 | x           a | sum | count
 1 | 3 | y          ---+-----+-------
 2 | 2 | z           1 |   5 |     2
 2 | 1 | a           2 |   6 |     3
 2 | 3 | b
```

User-defined Aggregates

SQL standard does not specify user-defined aggregates.

But PostgreSQL provides a mechanism for defining them.

To define a new aggregate, first need to supply:

- *BaseType* ... type of input values
- *StateType* ... type of intermediate states
- state mapping function: *sfunc(state,value)* → *newState*
- [optionally] an initial state value (defaults to null)
- [optionally] final function: *ffunc(state)* → *result*

... User-defined Aggregates

New aggregates defined using `CREATE AGGREGATE` statement:

```
CREATE AGGREGATE AggName(BaseType) (
    sfunc      = NewStateFunction,
```

```
    stype      = StateType,
    initcond   = InitialValue,
    finalfunc  = FinalResFunction,
    sortop     = OrderingOperator
);
```

- initcond (type *StateType*) is optional; defaults to NULL
- **finalfunc** is optional; defaults to identity function
- sortop is optional; needed for min/max-type aggregates

... User-defined Aggregates

Example: defining the count aggregate (roughly)

```
create aggregate myCount(anyelement) (    Element can be different types
    stype     = int,    -- the accumulator type
    initcond = 0,       -- initial accumulator value
    sfunc     = oneMore -- increment function
);
```

You should create a function called "oneMore"

```
create function
    oneMore(sum int, x anyelement) returns int
as $$
begin return sum + 1; end;
$$ language plpgsql;
```

... User-defined Aggregates

Example: sum2 sums two columns of integers

```
create type IntPair as (x int, y int);

create function
    AddPair(sum int, p IntPair) returns int
as $$
begin return p.x + p.y + sum; end;
$$ language plpgsql;
```

```
create aggregate sum2(IntPair) (
    stype     = int,
    initcond  = 0,
    sfunc     = AddPair
);
```

Exercise: Defining Aggregates

Define a `concat` aggregate that

- takes a column of string values
- returns a comma-separated string of values

Example:

```
select count(*), concat(name) from Employee;
-- returns e.g.
  count |        concat
 -------+---------------------
      4 | John,Jane,David,Phil
```

Use it to get a list of beers liked by each drinker.

Constraints

Constraints

So far, we have considered several kinds of constraints:

- attribute (column) constraints
- relation (table) constraints
- referential integrity constraints

Examples:

```
create table Employee (
    id       integer primary key,
    name     varchar(40),
    salary   real,
    age      integer check (age > 15),
    worksIn integer
              references Department(id),
    constraint PayOk check (salary > age*1000)
);
```

... Constraints

Column and table constraints ensure validity of one table.

RI constraints ensure connections between tables are valid.

However, specifying validity of entire database often requires constraints involving multiple tables.

Simple example (from banking domain):

```
for all Branches b
    b.assets == (select sum(acct.balance)
                 from   Accounts acct
                 where  acct.branch = b.location)
```

i.e. assets of a branch is sum of balances of accounts held at that branch

Assertions

Assertions are schema-level constraints

- typically involving multiple tables
- expressing a condition that must hold at all times
- need to be checked on each change to relevant tables
- if change would cause check to fail, reject change

SQL syntax for assertions:

```
CREATE ASSERTION name CHECK (condition)
```

The *condition* is expressed as "there are no violations in the database"

Implementation: ask a query to find all the violations; check for empty result

---

... Assertions

Example: #students in any UNSW course must be < 10000

```
create assertion ClassSizeConstraint check (
   not exists (
      select c.id from Courses c, CourseEnrolments e
      where  c.id = e.course
      group  by c.id having count(e.student) > 9999
   )
);
```

Needs to be checked after *every* change to either `Course` or `CourseEnrolment`

... Assertions

Example: assets of branch = sum of its account balances

```
create assertion AssetsCheck check (
   not exists (
      select branchName from Branches b
      where  b.assets <>
             (select sum(a.balance) from Accounts a
              where a.branch = b.location)
   )
);
```

Needs to be checked after *every* change to either `Branch` or `Account`

... Assertions

On each update, it is expensive

- to determine which assertions need to be checked
- to run the queries which check the assertions

A database with many assertions would be way too slow.

So, most RDBMSs do not implement general assertions.

Typically, *triggers* are provided as

- a lightweight mechanism for dealing with assertions
- a general event-based programming tool for databases

Triggers

Triggers

Triggers are

- procedures stored in the database
- activated in response to database events   (e.g. updates)

Examples of uses for triggers:

- maintaining summary data
- checking schema-level constraints (assertions) on update
- performing multi-table updates (to maintain assertions)

... Triggers

Triggers provide event-condition-action (ECA) programming:

- an *event* activates the trigger
- on activation, the trigger checks a *condition*
- if the condition holds, a procedure is executed (the *action*)

Some typical variations on this:

- execute the action before, after or instead of the triggering event
- can refer to both old and new values of updated tuples
- can limit updates to a particular set of attributes
- perform action: for each modified tuple, once for all modified tuples

---

... Triggers

SQL "standard" syntax for defining triggers:

```
CREATE TRIGGER TriggerName
{AFTER│BEFORE}  Event1 [ OR Event2 ... ]
[ FOR EACH ROW ]
ON TableName
[ WHEN ( Condition ) ]
Block of Procedural/SQL Code ;
```

Possible *Events* are `INSERT`, `DELETE`, `UPDATE`.

`FOR EACH ROW` clause ...

- if present, code is executed on each modified tuple
- if not present, code is executed once after all tuples are modified, just before changes are finally `COMMIT`ed

---

Trigger Semantics

Triggers can be activated `BEFORE` or `AFTER` the event.

If activated `BEFORE`, can affect the change that occurs:

- `NEW` contains "proposed" value of changed tuple
- modifying `NEW` causes a different value to be placed in DB

If activated `AFTER`, the effects of the event are visible:

- `NEW` contains the current value of the changed tuple
- `OLD` contains the previous value of the changed tuple
- constraint-checking has been done for `NEW`

Note: `OLD` does not exist for insertion; `NEW` does not exist for deletion.

... Trigger Semantics

Consider two triggers and an INSERT statement

```
create trigger X before insert on T Code1;
create trigger Y after insert on T Code2;
insert into T values (a,b,c,...);
```

Sequence of events:

- execute `Code1` for trigger `X`
- code has access to `(a,b,c,...)` via `NEW`
- code typically checks the values of `a,b,c,..`
- code can modify values of `a,b,c,..` in `NEW`
- DBMS does constraint checking as if `NEW` is inserted
- if fails any checking, abort insertion and rollback
- execute `Code2` for trigger `Y`
- code has access to final version of tuple via `NEW`
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: there is no `OLD` tuple for an `INSERT` trigger.

... Trigger Semantics

Consider two triggers and an UPDATE statement

```
create trigger X before update on T Code1;
create trigger Y after update on T Code2;
update T set b=j,c=k where a=m;
```

Sequence of events:

- execute `Code1` for trigger `X`
- code has access to current version of tuple via `OLD`
- code has access to updated version of tuple via `NEW`
- code typically checks new values of `b,c,..`

- code can modify values of `a,b,c,..` in `NEW`
- do constraint checking as if `NEW` has replaced `OLD`
- if fails any checking, abort update and rollback
- execute `Code2` for trigger `Y`
- code has access to final version of tuple via `NEW`
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: both `OLD` and `NEW` exist in UPDATE triggers.

... Trigger Semantics

Consider two triggers and an DELETE statement

```
create trigger X before delete on T Code1;
create trigger Y after delete on T Code2;
delete from T where a=m;
```

Sequence of events:

- execute `Code1` for trigger `X`
- code has access to `(a,b,c,...)` via `OLD`
- code typically checks the values of `a,b,c,..`
- DBMS does constraint checking as if `OLD` is removed
- if fails any checking, abort deletion (restore `OLD`)
- execute `Code2` for trigger `Y`
- code has access to about-to-be-deleted tuple via `OLD`
- code typically does final checking, or modifies other tables in database to ensure constraints are satisfied

Reminder: tuple `NEW` does not exist in DELETE triggers.

Triggers in PostgreSQL

PostgreSQL triggers provide a mechanism for

- `INSERT`, `DELETE` or `UPDATE` events
- to automatically activate PLpgSQL functions

Syntax for PostgreSQL trigger definition:

```
CREATE TRIGGER TriggerName
{AFTER│BEFORE}  Event1 [OR Event2 ...]
ON TableName
[ WHEN ( Condition ) ]
FOR EACH {ROW│STATEMENT}
EXECUTE PROCEDURE FunctionName(args...);
```

... Triggers in PostgreSQL

There is no restriction on what code can go in the function.

However a BEFORE function must contain one of:

RETURN old;   or   RETURN new;

depending on which version of the tuple is to be used.

If BEFORE trigger returns old, no change occurs.

If exception is raised in trigger function, no change occurs.

Produced: 23 Aug 2016